Adrian Salvador Ekomo

Computer Engineering Student — Data Engineering · Cloud & DevOps · SAP BW/BI

Tunis, Tunisia • adriansalvadorekomo@outlook.com • adriansalvadorekomo • Adrian Salvador Ekomo

Professional Summary

Computer Engineering student specializing in Data Engineering, building reliable platforms that turn data into insights. Hands-on with pipelines, warehouses, analytics, and ML workflows (Airflow, **dbt**, **Power BI**, Metabase, **MLflow**), plus professional DevOps experience (**OpenStack**, **Terraform**, automation). Focused on raw-to-trusted data for business decisions, learning SAP BW/BI. Seeking a Junior **Data Engineer** role.

Education

Esprit School of Engineering — Sept 2022 – June 2027, Engineering (PABI — Business Intelligence) in Computer Science – Tunis, Tunisia

- Specialized in Data Engineering, Business Intelligence, and Applied Machine Learning

Experience

DevOps Engineer · Internship (Remote), RIF — Rassemblement des Ingénieurs Francophones – La Marsa, Tunis July 2026 – Sept 2026

Built a secure, highly available Cloud infrastructure on **OpenStack** to host the Open edX learning platform (code).

- Provisioned VMs, networks, Floating IPs and Security Groups with **Terraform** (IaC) and configured them with idempotent **Ansible** roles (**Docker**, UFW, reverse proxy)
- Deployed containerized Open edX behind an HA Bastion pair (ProxyJump, validated Keepalived failover) served via Nginx/DuckDNS/HTTPS

Projects

Event Analytics Platform Jan 2026 – May 2026

End-to-end event analytics platform built during an academic project, featuring a star-schema data warehouse, **ETL** pipelines, ML models for forecasting and anomaly detection, NL-to-**SQL** query interface, and containerized microservice architecture.

- Built a star-schema warehouse in **PostgreSQL** that unified event, booking, and marketing data across 3 source systems for cross-functional analysis — organized around business questions, not source schemas. Used Talend for **ETL** transformations and **Apache Airflow** for reliable ingestion
- Developed ML models — Prophet for demand forecasting (handles event seasonality naturally) with MAPE under 15%, statsmodels for anomaly detection, and scikit-learn for classification — tracked through **MLflow** and exposed via a natural-language interface so anyone could explore the data
- Containerized 11 microservices with **Docker** so each component stayed deployable on its own, set up monitoring to catch issues early, automated ML retraining to keep models accurate as data patterns shifted, and reduced manual deployment steps and recovery time
- View code on GitHub

Environmental Monitoring AI Mar 2026 – Mar 2026

AI-powered environmental monitoring platform developed during a national innovation hackathon, integrating satellite image segmentation, time-series forecasting, air quality prediction, and a multilingual chatbot interface.

- Built satellite segmentation models with **PyTorch** on multispectral imagery to classify soil contamination and detect water turbidity — ~85% accuracy on soil classification, giving field teams real-time environmental data they could act on
- Used Airbyte to sync sensor data from 20+ field devices into a central database — eliminating manual data wrangling — then built time-series models with scikit-learn to forecast contamination trends and air quality with RMSE under 20%, turning noisy measurements into clear environmental insights
- View code on GitHub

DataOps Data Platform Apr 2026 – present

Data platform project integrating version-controlled **dbt** data models for analytics, ML models for business forecasting, and a **RAG**-based natural-language query assistant. In progress.

- Used Airbyte to sync CRM, sales, and support data (thousands of records/day) directly into the warehouse — cutting data latency from hours to <5 minutes — then set up **dbt** models running transformations inside the warehouse, keeping the pipeline simpler by doing the heavy work where data already lives. Used **Apache Airflow** to orchestrate model runs and built **PyTorch** forecasting models for sales and churn with **MLflow** tracking
- Built a **RAG**-based assistant that let business teams query **PostgreSQL** data in plain English — handling ~70% of ad-hoc queries without engineering tickets, so they could get answers faster
- View code on GitHub

Skills

Data Engineering: SQL, **PostgreSQL**, Talend, Airbyte, **Apache Airflow**, **ETL**/ELT, Data Warehousing, Data Modeling, Dimensional Modeling

Analytics & BI Engineering: dbt, **Power BI**, Metabase, SAP BW (learning), Prometheus, Grafana

ML/DL Model Integration & MLOps: Scikit-learn, statsmodels, **PyTorch**, **MLflow**, **RAG**, NLP-to-**SQL**, Time-Series Forecasting, Clustering, Anomaly Detection

Backend & DevOps: FastAPI, **REST** APIs, **Docker**, **Linux**, **Terraform**, **Ansible**, **OpenStack**, Nginx, n8n

Programming Languages: Python, R, Java, JavaScript/TypeScript, C

Languages

Spanish: Native — Full professional proficiency

English: B2 — Read ML/AI papers, write technical docs

French: B2 — Read technical papers, write documentation

German: A1 — Basic

Certifications

Python Data Associate — DataCamp (Sept 2026)

Data Engineer Associate — DataCamp (Sept 2026)

SQL Associate — DataCamp (Sept 2026)