

Adrian Salvador Ekomo

Computer Engineering Student — Data Engineering · Cloud & DevOps · SAP BW/BI

📍 Tunis, Tunisia • ✉️ adriansalvadorekomo@outlook.com • 🌐 [adriansalvadorekomo](#) • 📺 Adrian Salvador Ekomo

Professional Summary

Computer Engineering student focused on becoming a data engineer: pipelines, warehouses, and analytics with **dbt** + Airflow. Differentiator: Cloud/DevOps with **Terraform** + **OpenStack** from a production-style internship. Complementary ML (**MLflow**) and BI (**Power BI**, Metabase); learning SAP BW/BI. Seeking a Junior **Data Engineer** role.

Education

Esprit School of Engineering — Sept 2022 – June 2027, Engineering (PABI — Business Intelligence) in Computer Science – Tunis, Tunisia

- Specialized in Data Engineering, Business Intelligence, and Applied Machine Learning

Experience

DevOps Engineer · Internship (Remote), RIF — Rassemblement des Ingénieurs Francophones – La Marsa, Tunis July 2026 – Sept 2026

My scope: **Terraform/OpenStack** provisioning, **Ansible** configuration, containerized Open edX deployment, and HA Bastion validation ([code](#)).

- Provisioned VMs, networks, Floating IPs and Security Groups with **Terraform** configured **Docker**, UFW and reverse proxy with idempotent **Ansible** roles
- Validated PRIMARY→BACKUP failover in ~5s on one Floating IP; served Open edX via Nginx/DuckDNS/HTTPS

Projects

Event Analytics Platform

Jan 2026 – May 2026

Academic team project; my scope: integration, **Docker**/nginx delivery, and chatbot/API fixes. End-to-end event analytics platform with star-schema warehouse, **ETL**, ML forecasting/anomaly detection, NL-to-**SQL**, and containerized microservices.

- Built a star-schema warehouse in **PostgreSQL** that unified event, booking, and marketing data across 3 source systems for cross-functional analysis — organized around business questions, not source schemas. Used Talend for **ETL** transformations and **Apache Airflow** for reliable ingestion
- Developed ML models — Prophet for demand forecasting (handles event seasonality naturally) with MAPE under 15%, statsmodels for anomaly detection, and scikit-learn for classification — tracked through **MLflow** and exposed via a natural-language interface so anyone could explore the data
- Containerized 11 microservices with **Docker** behind an nginx proxy, fixed chatbot/API integration, and automated monitoring plus ML retraining
- [View code on GitHub](#)

Environmental Monitoring AI

Mar 2026 – Mar 2026

AI-powered environmental monitoring platform developed during a national innovation hackathon, integrating satellite image segmentation, time-series forecasting, air quality prediction, and a multilingual chatbot interface.

- Built satellite segmentation models with **PyTorch** on multispectral imagery to classify soil contamination and detect water turbidity — ~85% accuracy on soil classification, giving field teams real-time environmental data they could act on
- Used Airbyte to sync sensor data from 20+ field devices into a central database — eliminating manual data wrangling — then built time-series models with scikit-learn to forecast contamination trends and air quality with RMSE under 20%, turning noisy measurements into clear environmental insights
- [View code on GitHub](#)

DataOps Data Platform

Apr 2026 – present

Personal in-progress DataOps build: version-controlled **dbt** models, ML business forecasting, and **RAG**-based NL query assistant.

- Used Airbyte to sync CRM, sales, and support data (thousands of records/day) directly into the warehouse — cutting data latency from hours to <5 minutes — then set up **dbt** models running transformations inside the warehouse, keeping the pipeline simpler by doing the heavy work where data already lives. Used **Apache Airflow** to orchestrate model runs and built **PyTorch** forecasting models for sales and churn with **MLflow** tracking
- Built a **RAG**-based assistant that let business teams query **PostgreSQL** data in plain English — handling ~70% of ad-hoc queries without engineering tickets, so they could get answers faster
- [View code on GitHub](#)

Skills

Data Engineering: **SQL**, **PostgreSQL**, Talend, Airbyte, **Apache Airflow**, **ETL/ELT**, Data Warehousing, Data Modeling, Dimensional Modeling

Cloud & DevOps: **Terraform**, **OpenStack**, **Ansible**, **Docker**, **Linux**, **FastAPI**, **REST** APIs, Nginx, n8n

Analytics & BI Engineering: **dbt**, **Power BI**, Metabase, SAP BW (learning), Prometheus, Grafana

ML/DL Model Integration & MLOps: Scikit-learn, statsmodels, **PyTorch**, **MLflow**, **RAG**, NLP-to-**SQL**, Time-Series Forecasting, Clustering, Anomaly Detection

Programming Languages: **Python**, R, Java, JavaScript/TypeScript, C

Languages

Spanish: Native — Full professional proficiency

English: B2 — Read ML/AI papers, write technical docs

French: B2 — Read technical papers, write documentation

German: A1 — Basic

Certifications

Python Data Associate — DataCamp (Sept 2026)

Data Engineer Associate — DataCamp (Sept 2026)

SQL Associate — DataCamp (Sept 2026)

Prompt Engineering for Everyone — Cognitive Class / IBM (July 2026)